

ADVANCED

overall tree
= binary tree of mini trees

mini trees

micro trees

actual nodes

$\frac{1}{7} \lg n$ nodes

$\lg n$ nodes

n nodes

DATA STRUCTURES

2

Adaptive Trees

Advanced Data Structures · Summer 2026

Prof. Dr. Sebastian Wild


Outline

2 Adaptive Trees

- 2.1 Move-to-Root Heuristic
- 2.2 Move-to-Root Analysis
- 2.3 Splay Trees
- 2.4 Analysis of Splay Trees
- 2.5 Biased Search Trees
- 2.6 Excursion: Online Algorithms
- 2.7 Dynamic Optimality
- 2.8 The Geometric View of BSTs
- 2.9 Deferred Data Structures
- 2.10 Lazy Search Trees
- 2.11 Lazy Search Tree Operations



Self-Adjusting Data Structures

Idea: set of rules that makes data structure automatically adapt to use case



2.1 Move-to-Root Heuristic

Rotations



Move-to-root

Upon an access to an element, rotateParent until accessed element is at root!



Allen, Munro: *Self-Organizing Binary Search Trees*, JACM 1978

Inspired by **Move-To-Front heuristic** of self-adjusting linked lists

- ▶ cost via MTF is at most 2 times cost of optimal **static ordering**
- ▶ MTF is 2-competitive as an online algorithm
(at most 2 times the cost of *any* algorithm)
- ▶ (both true under “basic cost model” only)
- ▶ with resource augmentation, LRU in paging very successful

Maybe, Move-to-root is similarly successful!

2.2 Move-to-Root Analysis

Iid Accesses

Iid Model: Suppose the accesses to keys $[1..n]$ in BST are *i.i.d. randomly* drawn with probabilities $p_1, \dots, p_n$.

↪ Expected search cost in a BST T given by $\sum_{i=1}^n p_i \cdot \text{depth}_T(\textcircled{i})$

► If p_i known up front, can compute **static optimal BST** via dynamic programming

↪ Cost of this optimal tree C_{OPT} can be bounded in terms of the entropy

Lemma 2.1 (Bayer 1975)

$$\mathcal{H} - \lg \mathcal{H} - \lg e + 1 \leq C_{OPT} \leq \mathcal{H} + 1 \quad \text{for} \quad \mathcal{H} = \sum_{i=1}^n p_i \lg(1/p_i)$$



Bayer: *Improved Bounds on the Costs of Optimal and Balanced Binary Search Trees*, M.Sc. Thesis, MIT 1975

Note: C_{OPT} is for successful search / depth of internal nodes

The case of accesses to leaves only, with access to $\boxed{j}$ w/p $q_j, j = 0, \dots, n$ even cleaner:

$$\mathcal{H} \leq \sum_{j=0}^n q_j \cdot (1 + \text{depth}_T(\boxed{j})) \leq \mathcal{H} + 2$$

Move-to-Root on iid Accesses

Lemma 2.2 (MTR ancestors)

Let T be a tree maintained by the move-to-root heuristic and $i < j$ two keys.

Then (i) is an **ancestor** of (j) iff the most recent request for i came after the most recent requests of any of $i + 1, \dots, j$.

Similarly, (j) ancestor of (i) iff j requested more recently than any of $i, \dots, j - 1$. ◀

Recall $A_j^i = [(i) \text{ ancestor of } (j)]$

Corollary 2.3

Under the iid accesses model and MTR heuristic, $\mathbb{P}[A_j^i] = \begin{cases} \frac{p_i}{p_i + \dots + p_j} & i < j \\ \frac{p_i}{p_j + \dots + p_i} & i > j \end{cases}$ ◀

Move-to-Root on iid Accesses – Analysis

Theorem 2.4 (MTR cost)

Let C_{MTR} be the expected search cost under the i.i.d. model in a tree T maintained by MTR.

$$C_{MTR} = 1 + 2 \sum_{1 \leq i < j \leq n} \frac{p_i p_j}{p_i + \dots + p_j} \leq 1 + 2 \ln 2 \cdot \mathcal{H}$$

Move-to-Root has close to optimal search cost in the i.i.d. model!

Proof:

$$\begin{aligned} C_{MTR} &= \sum_{i=1}^n p_i \cdot (1 + \text{depth}(\textcircled{i})) = \sum_{i=1}^n p_i \cdot \left(1 + \sum_{j \neq i} \mathbb{E}[A_i^j] \right) \\ &= 1 + \sum_{i=1}^n \sum_{j \neq i} p_i \mathbb{E}[A_i^j] \\ &= 1 + 2 \sum_{1 \leq i < j \leq n} \frac{p_i p_j}{p_i + \dots + p_j} \end{aligned}$$

This proves the first part.

Move-to-Root on iid Accesses – Analysis

Proof (cont.):

Mehlhorn's trick: $\delta_{i,j} := \frac{p_i + \dots + p_j}{p_i} \rightsquigarrow 1 = \delta_{i,i} \leq \delta_{i,i+1} \leq \dots \leq \delta_{i,n} \leq \frac{1}{p_i}$

$$C_{MTR} = 1 + 2 \sum_{i=1}^{n-1} p_i \sum_{j=i+1}^n \frac{p_j}{p_i + \dots + p_j}$$

$$= 1 + 2 \sum_{i=1}^{n-1} p_i \sum_{j=i+1}^n \frac{\delta_{i,j} - \delta_{i,j-1}}{\delta_{i,j}}$$

$$a \leq b \implies (b-a) \cdot \frac{1}{b} = \int_a^b \frac{1}{b} dx < \int_a^b \frac{1}{x} dx$$

$$< 1 + 2 \sum_{i=1}^{n-1} p_i \sum_{j=i+1}^n \int_{x=\delta_{i,j-1}}^{\delta_{i,j}} \frac{1}{x} dx$$

$$= 1 + 2 \sum_{i=1}^{n-1} p_i \int_{x=\delta_{i,i}}^{\delta_{i,n}} \frac{1}{x} dx$$

$$\int_{x=a}^b \frac{1}{x} dx = \ln(b) - \ln(a) = \ln(b/a)$$

$$= 1 + 2 \sum_{i=1}^{n-1} p_i \ln(\delta_{i,n}/\delta_{i,i}) \leq 1 + 2 \ln 2 \sum_{i=1}^{n-1} p_i \lg(1/p_i) = 1 + 2 \ln 2 \cdot \mathcal{H}$$

Unbalanced BSTs under i.i.d. model

Interestingly, the same cost result from unbalanced BSTs!

- ▶ Suppose we repeatedly draw x i.i.d. from $[1..n]$ w/p $p_1, \dots, p_n$
- ▶ We insert x into initially empty unbalanced BST T
- ▶ **If x is already contained in T , the insertion has no effect**
- ▶ Repeat until saturation (i. e., until we have seen every value $x \in [1..n]$)

Since insertions are i.i.d., probability that i is inserted first among $[i..j]$ is $\frac{p_i}{p_i + \dots + p_j}$

$$\rightsquigarrow \mathbb{P}[A_j^i] = \begin{cases} \frac{p_i}{p_i + \dots + p_j} & i < j \\ \frac{p_i}{p_j + \dots + p_i} & i > j \end{cases}$$

$$\rightsquigarrow C_T = 1 + 2 \sum_{1 \leq i < j \leq n} \frac{p_i p_j}{p_i + \dots + p_j} \leq 1 + 2 \ln 2 \cdot \mathcal{H}$$

Both unbalanced BST and Move-to-Root behave well under i.i.d. model.

Move-to-Root's Blind Spot

Problem: *Can't rely on accesses to be random.*

Theorem 2.5 (Move-to-Root Bad Case)

There exists a sequence of n requests, such that starting from any tree T , the amortized access cost over the course of the n requests is $\Omega(n)$.

Proof:

see exercises

2.3 Splay Trees

Splay Trees

Splay trees by now the most widely known self-adjusting data structure.

👍 reasonably simple and fast

👍 simple to implement

👍 remarkable theoretical properties (*access theorems*) $\rightsquigarrow$ later


👎 many writes and pointer chases

📖 Sleater, Tarjan: *Self-Adjusting Binary Search Trees*, JACM 1985


Move-to-root reloaded

What's wrong in Move-to-front?

Let's look at the possible cases of parent and grandparent




Using simple `rotateParent` calls, the subtree we came from (B or C) is lifted up!



Here, simple `rotateParent` calls do **not** lift C up!

The Fix


*In the case that both parents are right children, do a **double rotation***



Now both C and D are lifted up!

Splay


Aggressive restructuring allows to base all operations on single primitive: splay



```
1 def splay(T, x):
2     while x ≠ T.root
3         p = parent(x)
4         if p == T.root
5             rotateParent(x) # Zig Case
6         else:
7             g = parent(p)
8             if p == g.left ∧ x == p.left or p == g.right ∧ x == p.right:
9                 doubleRotateParent(x) # Zig-Zig Case
10            else:
11                rotateParent(x) # Zig-Zag Case
12                rotateParent(x)
```


Splay – Insert

1. $\text{splay}(x, T) \rightsquigarrow$ root y and $T_\ell \leq y \leq T_r$ (key order)
2. If $x = y$ (key already stored) return.
3. Otherwise, w.l.o.g. $x < y$, can simply make x new root.



Splay – Delete

1. $\text{splay}(x, T) \rightsquigarrow$ root x (or not part of tree and we're done)
2. Remove x , left with T_ℓ and T_r
3. $\text{splay}(x, T_\ell) \rightsquigarrow$ root $y = \max(T_\ell)$ and $y.\text{right} = \text{null}$
4. $y.\text{right} := T_r$



Note: We've effectively built **join** $\rightsquigarrow$ naturally supported via splay.

split even more directly supported via splay: $\text{split}(T, x)$: $\text{splay}(T, x)$, return $(x.\text{left}, x)$

2.4 Analysis of Splay Trees

Rotation Cost

Note: In all operations, we follow a search path and splay.

(For delete twice)

splay walks back up search path

↪ use $k = \Theta(d)$ rotations for splaying node at depth d .

↪ cost is dominated by $k = \text{\#rotations in splay}$ calls.

Clearly, individual operations could be expensive (use $k = \Omega(n)$ rotations).

↪ Can only hope for an *amortized* bound on k

We define a potential function $\Phi = \Phi(T)$ such that

Goal: If $\text{splay}(x)$ uses k rotations, then

$$k + \Phi' - \Phi \leq 1 + 3 \cdot (r'(x) - r(x))$$

$\Delta\Phi = \text{change in potential}$

change in "rank" of x

↪ Released potential can pay for actual cost k

Splay Tree Potential

We choose a general method that will give several cool results.

- ▶ For that, we allow each node x to have a *weight* $w(x)$

- ▶ The *size* of a node x is defined as
$$s(x) := \sum_{\substack{v \in T: \\ x \text{ ancestor of } v}} w(v)$$

- ▶ The *rank* of node x is
$$r(x) := \lg(s(x))$$

- ▶ The *potential* of a tree is
$$\Phi = \sum_{v \in T} r(v)$$

- ▶ We always denote with Φ, r, s the quantities **before** an operation and with Φ', r', s' the quantities **after** an operation

The Access Lemma

```
1 def splay(T, x):
2     while x ≠ T.root
3         p = parent(x)
4         if p == T.root
5             rotateParent(x) # Zig Case
6         else:
7             g = parent(p)
8             if p == g.left ∧ x == p.left or p == g.right ∧ x == p.right:
9                 doubleRotateParent(x) # Zig-Zig Case
10            else:
11                rotateParent(x) # Zig-Zag Case
12                rotateParent(x)
```

$$\Phi = \sum_{v \in T} r(v)$$

$$r(x) := \lg(s(x))$$

$$s(x) := \sum_{\substack{v \in T: \\ x \text{ ancestor of } v}} w(v)$$

Lemma 2.6 (Access Lemma)

For **any** weakly positive node weights w , the amortized cost

- (a) for a **Zig** is $\leq 1 + 3(r'(x) - r(x))$
- (b) for a **Zig-Zig** is $\leq 3(r'(x) - r(x))$
- (c) for a **Zig-Zag** is $\leq 3(r'(x) - r(x))$



Access Lemma – Proof

Splay Cost

Corollary 2.7

The amortized cost to splay a tree with root t at node x is at most

$$3(r(t) - r(x)) + 1 = O\left(\log\left(\frac{s(t)}{s(x)}\right)\right)$$

Proof:

Splay does rotations involving the same node x until it is the root.

~> The rank differences telescope.

~> Result follow from Lemma 2.6.

Lemma 2.8 (Potential drop)

If weights $w(x)$ are fixed, the potential drop over a sequence of m splay operations is

$$\Phi_0 - \Phi_m \leq \sum_{x \in T} \lg\left(\frac{W}{w(x)}\right) \quad \text{where} \quad W = \sum_{x \in T} w(x)$$

Proof:

Always have $\Phi \leq n \lg W$ and $\Phi \geq \sum_{x \in T} \lg w(x)$ (by definition)

So $\Phi_m \leq n \lg W$ and $\Phi_0 \geq \sum \lg w(x)$.

Splay Tree – Results

Theorem 2.9 (Balance Theorem)

Consider a splay tree containing n keys and an arbitrary sequence of m accesses to these keys. The total access time is $O(n \log n + m \log n)$. ◀

Splay Tree – Results [2]

Fix an access sequence $X = x_1, \dots, x_m$ on the n nodes in splay tree T , and denote by $p_i = |X|_i/m$ the relative frequency of i in X .

Theorem 2.10 (Static Optimality)

Assume access sequence X contains all n nodes (i. e., $p_i > 0$ for all i).

A splay tree serving access X incurs total cost $O(m(\mathcal{H} + 1))$ for $\mathcal{H} = \sum_{i=1}^n p_i \lg(1/p_i)$. ◀

Splay Tree – Results [3]

Theorem 2.11 (Static Finger Theorem)

For any fixed node f (the “finger”), the total cost of access sequence X is

$$O\left(m + n \log n + \sum_{j=1}^m \log(1 + |x - y|)\right)$$

Here $|x - y|$ is the *rank distance*, i. e., the number of nodes $z \in T$ with $x \leq z < y$ (or $y \leq z < x$)

Splay Tree – Results [4]

For access sequence $X = x_1, \dots, x_m$ define $t(j)$ as the **number of different** items accessed since the **previous** access to the same node x_j (or since beginning)

Theorem 2.12 (Working Set Theorem)

The total cost of access sequence X is $O\left(m + n \log n + \sum_{j=1}^m \log(1 + t(j))\right)$



What makes Splay tick?

When you rotate constantly, it is easy to lose orientation in the forest . . .

A global view of Splay




Fig. 2: A global view of splay trees. The transformation from the left to the middle illustrates rotate-to-root. The transformation from the left to the right illustrates splay trees.

Splay: Splay extends rotate-to-root: Let $s = v_0, v_1, \dots, v_k$ be the reversed search path. We view splaying as a two step process, see Figure 2. We first make s the root and split the search path into two paths, the path of elements smaller than s and the path of elements larger than s . If v_{2i+1} and v_{2i+2} are on the same side of s , we rotate them, i.e., we remove v_{2i+2} from the path and make it a child of v_{2i+1} .

Proposition 17. *The above description of splay is equivalent to the Sleator-Tarjan description.*



Chalermsook, Goswami, Kozma, Mehlhorn, Saranurak: *Self-Adjusting Binary Search Trees: What Makes Them Tick?*, ESA 2015


Intuition:

- ▶ Move-to-root = treap with newly accessed element gets *new maximal priority*
- ▶ Splay = that, plus *demotion of even index if previous on same side*

Anything magic about Splay?

Same paper shows: whole **class** of rearrangements satisfy access lemma

Characteristic quantities of the **search path** and the **after-tree**:



- length of search path: $|P|$ 12
- number of zigzags: z 4
- number of leaves: ℓ 5
- max left-depth or right-depth: d 3

Theorem. If, for every search:

- (1) max left/right depth: $d = O(1)$,
- (2) number of leaves+zigzags: $\ell + z = \Omega(|P|)$

$\Rightarrow$ Algorithm shares many known good properties of Splay.

2.5 Biased Search Trees

Updates in Splay Trees

We've already seen: insert, delete, split, join all implemented on top of splay

- ▶ suffices to analyze these splay calls.
- ▶ but: potential so far only defined for fixed tree!

Dynamic Splay Tree Potential

- ▶ **general state:** collection $\mathcal{T}$ of splay trees (initially empty)
- ▶ all items ever inserted exist from the beginning in the “item ether” $\mathcal{E}$ (outside any tree)
deleted items go back to the ether

$\rightsquigarrow$ universe U of items is fixed

- ▶ assume every item in at most one tree at any time

$$\rightsquigarrow \Phi = \sum_{T \in \mathcal{T}} \sum_{x \in T} \lg(s(x)) + \sum_{x \in \mathcal{E}} \lg(w(x))$$

Notation: For $x \in U$ and T clear from context, abbreviate
 $x_- = \text{predecessor}_T(x)$ and $x_+ = \text{successor}_T(x)$

Analysis of Updates

Theorem 2.13 (Splay Tree Update Lemma)

For any assignment of weights $w : U \rightarrow \mathbb{R}_{>0}$, the amortized operation costs are upper bounded by

- (a) $\text{access}_T(x)$ where $x \in T$: $3 \lg \left(\frac{W}{w(x)} \right) + 1$ for $W = \sum_{x \in T} w(x)$
- (b) $\text{access}_T(x)$ where $x \notin T$: $3 \lg \left(\max \left\{ \frac{W}{w(x_-)}, \frac{W}{w(x_+)} \right\} \right) + 1$
- (c) $\text{join}(T_1, T_2)$: $3 \lg \left(\frac{W}{w(x)} \right) + O(1)$ for $x = \max(T_1)$, $W = \sum_{x \in T_1} w(x) + \sum_{x \in T_2} w(x)$
- (d) $\text{split}_T(x)$ where $x \in T$: $3 \lg \left(\frac{W}{w(x)} \right) + O(1)$ for $W = \sum_{x \in T} w(x)$
- (e) $\text{split}_T(x)$ where $x \notin T$: $3 \lg \left(\max \left\{ \frac{W}{w(x_-)}, \frac{W}{w(x_+)} \right\} \right) + O(1)$
- (f) $\text{insert}_T(x)$: $3 \lg \left(\max \left\{ \frac{W - w(x)}{w(x_-)}, \frac{W - w(x)}{w(x_+)} \right\} \right) + \lg \left(\frac{W}{w(x)} \right) + O(1)$
- (g) $\text{delete}_T(x)$: $3 \lg \left(\frac{W}{w(x)} \right) + 3 \lg \left(\frac{W - w(x)}{w(x_-)} \right) + O(1)$



Analysis of Updates – Proof

Biased Search Trees

Biased search trees are a refinement of sorted dictionaries where

- ▶ elements have *known* weights $w(x)$ (that the data structure can use)
- ▶ supports **biased** search costs (as in Theorem 2.14 when using these weights)
- ▶ there is an explicit **changeWeight_T**(x, δ) operation that sets $w(x) := w(x) + \delta$
- ▶ prior to Splay Trees solved via *globally-biased 2, b trees*
 -  Bent, Sleator, Tarjan: *Biased search trees*, SICOMP 1985
 - ↪ achieves all operations in **worst case** biased time
 - ▶ rather intricate invariants and higher overhead
 - ▶ biased access times also achievable via randomization: *treaps, zip-zip trees*

Weight Changes

Lemma 2.14 (Weight Changes)

For any assignment of weights $w : U \rightarrow \mathbb{R}_{>0}$, the amortized operation costs are upper bounded by

- (a) $\text{changeWeight}_T(\text{root}(T), \delta)$ where $\delta > 0$: $\lg\left(1 + \frac{\delta}{W}\right)$ for $W = \sum_{x \in T} w(x)$
- (b) $\text{changeWeight}_T(\text{root}(T), \delta)$ where $\delta < 0$: 0

Note: For changeWeight on arbitrary item x , first $\text{splay}(x)$ at cost $3 \lg\left(\frac{W}{w(x)}\right) + 1$.

2.6 Excursion: Online Algorithms

Online Problems

Online Problem

- ▶ algorithmic problem where input is only revealed over time
- ▶ and decisions/outputs need to be made along the way
- ▶ Typical example: stock trading, many scheduling problems
- ▶ default mode of thinking for **data structures!**
 - ▶ usually must answer one query before we receive the next

Opposite: Offline Problems

- ▶ get entire sequence of requests up front
- ⇒ can work out solution with foresight

Elevator (a.k.a. Ski Rental) Problem

The Elevator Problem*

*Rent skis or buy when not known how many days of snow left.

- ▶ Suppose you arrive in a building and need to go k floors up.
- ▶ You can either take the stairs, which takes you 1 min per floor $\rightsquigarrow k$ min.
- ▶ Or take the elevator, which takes only 1 min for the k floors $\dots$ *once it has arrived!*
- ▶ The elevator arrives after w min, where you know $w \in [0..B)$ $\rightsquigarrow w + 1$ min
- ▶ Assume B is a finite bound, but $B \gg k$.

What should you do (to minimize time)?

- ▶ Just walk?
- ▶ Wait for 2 min, then walk?
- ▶ Wait as long as it takes?

Impossible to tell what is better without some assumption on w , e. g., distribution assumption.

Elewaiting as Sequence of Requests

- ▶ formally, instance of online problem is sequence of requests $x = x_1, \dots, x_T$
- ▶ our algorithm $\mathcal{A}$ must produce output after each request
At time t , output $y_t = \mathcal{A}(x_1, \dots, x_t)$
- ▶ overall cost on instance depends on sequence of outputs
 $A(x) = \text{cost}(\mathcal{A}(x_1), \dots, \mathcal{A}(x_1, \dots, x_T))$

↪ Elevator Problem

- ▶ $x_t = [\text{elevator arrives at time } t], \quad T = B$
- ▶ $y_t = [\text{take stairs at time } t]$

Competitive Analysis

Elegant alternative to (potentially unrealistic) random models: *competitive analysis*

- ▶ Suppose $OPT(x)$ is the (cost of the) **optimal offline solution** for x
Note: OPT knows the future (entire access sequence x)
- ▶ For the elevator problem, OPT would wait if $w \leq k - 1$ and walk right away otherwise.
 $\rightsquigarrow OPT(w) = \min\{w + 1, k\}$

Definition 2.15 (Competitive ratio)

The *competitive ratio* of an online algorithm $\mathcal{A}$ (for size n) is

$$c = \max_x \frac{A(x)}{OPT(x)} \quad \text{where the maximum is taken over all instances of size } n.$$

We say that $\mathcal{A}$ is *c-competitive*.

- ▶ competitive ratio of “Just walk” is $\max_{w \in [0..B)} \frac{k}{\min\{k, w + 1\}} = k$
- ▶ “Always wait” has $\max_{w \in [0..B)} \frac{w + 1}{\min\{k, w + 1\}} = \frac{B}{k}$

Competitive Elewaiting

Can do much better by hedging out bets a bit: $\mathcal{A}_t$ waits t min then walks

$$\blacktriangleright A_t(w) = \begin{cases} w + 1 & w \leq t \\ t + k & w > t \end{cases}$$

$\rightsquigarrow$ Optimal choice for t

$$\arg \min_t \max_{w \in [0..B)} \frac{A_t(w)}{OPT(w)} = k$$

$\rightsquigarrow \mathcal{A}_k$ is 2-competitive

Further Famous Online-Algorithms

- ▶ Move-to-Front is 4-competitive for the list-update problem assuming only local swaps are allowed
- ▶ LRU is k -competitive in online paging (maintaining a cache of size k)
- ▶ Online bin packing (which box to put new item into, repacking not allowed)

2.7 Dynamic Optimality

Adaptive Trees as Online Algorithm

- ▶ Serving unknown access sequence to nodes of a BST is an online problem *par excellence*
Note: Focus here on access-only (no insertions & deletions).
- ▶ Splay trees are an online algorithm for this problem!
explicitly designed to do well in spite of not knowing the future

Compare to what

- ▶ *OPT* gets entire access sequence in advance
- ▶ but need to fix what it is allowed to do! when is an access counted as served?

The BST Model of Computation

- ▶ At any point in time, maintain a BST over $[1..n]$, and a *finger* (pointer) at one node.
Initial tree can be chosen by algorithm. (doesn't matter if access sequence long)
- ▶ Each step of the “computation” is one of the following 5 operations:
 - ▶ **move** finger from the current node to **left child**, **right child**, or **parent** (assuming they exist)
 - ▶ **rotate** the parent edge of the current node (assuming not at root)
 - ▶ **find**: report the current node as the sought value, reset finger to root of tree

↪ *OPT finds **shortest** valid sequence of operations serving accesses*

The Splay Tree Conjecture

Since Splay trees have all the strong adaptive properties, Sleator and Tarjan immediately conjectured:

CONJECTURE 1 (DYNAMIC OPTIMALITY CONJECTURE). *Consider any sequence of successful accesses on an n -node search tree. Let A be any algorithm that carries out each access by traversing the path from the root to the node containing the accessed item, at a cost of one plus the depth of the node containing the item, and that between accesses performs an arbitrary number of rotations anywhere in the tree, at a cost of one per rotation. Then the total time to perform all the accesses by splaying is no more than $O(n)$ plus a constant times the time required by algorithm A .*

(paper predates the term competitive ratio ...)

In today's words: **CONJECTURE 1:** Splay is $O(1)$ -competitive.

We're far from settling that conjecture ...

OPEN: *Is any online BST $o(\log \log n)$ -competitive?*

OPEN: *Are Splay trees $o(\log n)$ -competitive?*

► Clearly, $OPT(x) = \Omega(m)$ for $m = |x| \rightsquigarrow$ any static balanced tree is $O(\log n)$ -competitive

Rotation Distance & Initial Trees

Lemma 2.16

Using at most $2n - 2$ rotations, we can transform any BST on $[1..n]$ into any other BST on the same keys. ◀

Note, actually $2n - 6$ suffice.

Simplified BST Model

equivalent to original model up to constant factors, but easier to design algorithms

- ▶ rearrange “top part” (touched part) of tree (connected subgraph of tree containing root)
 - ▶ cost = number of nodes in top part
 - ▶ (isolated rotations in tree can be delayed)
 - ▶ **OPEN:** *Is restricting changes to **search path** w.l.o.g. in BST algorithms?*
 (“strict model”)
 - ▶ all sensible online algorithms work that way ...
- ▶ access only at root
 - ▶ at constant-factor overhead, can rotate up and back down again.

Facts on OPT

- ▶ can be computed in exponential time (already not trivial)
- ▶ exact *OPT* likely NP-hard

If each access is a set of items that the algorithm may order as convenient, the problem is known to be NP-complete

Lower Bounds for OPT

- ▶ Most access sequences must be “hard”, i. e., require $\Omega(m \log n)$ costs even for OPT
 - ▶ execution of BST algorithm **encodes** access sequence (must serve each differently!)
 - ▶ number of access sequences is n^m
 - ▶ each step of a BST algorithm is one of $O(1)$ choices
- ▶ The bitwise reversal sequence R_n needs $OPT(R_n) = \Omega(n \log n)$ cost.



Wilber. *Lower Bounds for Accessing Binary Search Trees with Rotations*, SICOMP 1989

- ▶ defined for $n = 2^k$
- ▶ $R_n = (\text{rev}_k(0), \text{rev}_k(1), \text{rev}_k(2), \dots, \text{rev}_k(n-1))$, $\text{rev}_k(i) = i$ in binary written in reverse.
 $R_{16} = (0, 8, 4, 12, 2, 10, 6, 14, 1, 9, 5, 13, 3, 11, 7, 15)$
- ▶ Wilber proved further lower bounds on OPT

Tango Trees

- ▶ Online BST algorithm specifically designed to stay close to one of Wilber's lower bounds
details skipped
- ▶ achieves a competitive ratio of $O(\log \log n)$
- ▶ may use $\Omega(\log n \log \log n)$ amortized cost $\rightsquigarrow$ not $o(\log \log n)$ -competitive



Demaine, Harmon, Iacono, Pătraşcu: *Dynamic Optimality – Almost*, SICOMP 2007

Greedy Future

Natural **offline** idea: find element, then rearrange search path as treap with next-future-access time (to element or subtrees) as priority

- ▶ Of course an **offline** method since it uses the future access times
- ▶ Long conjectured to be close to optimal ($+m$ over OPT ?)
- ▶ Now known to be at best a 2-approximation, and has cost $OPT + \Omega(m \log \log n)$ on some access sequences



Sadeh, Kaplan: *Dynamic Binary Search Trees: Improved Lower Bounds for the Greedy-Future Algorithm*, STACS 2023

- ▶ *May still be an offline $O(1)$ -approximation to OPT*
OPEN: Is GreedyFuture a $o(\log n)$ -approximation?

2.8 The Geometric View of BSTs

Different Way of Looking at BSTs

Beautiful reformulation of problem

 Demaine, Harmon, Iacono, Kane, Pătraşcu: *The Geometry of Binary Search Trees*, SODA 2009

- ▶ inspired a new contender for a dynamically optimal BST algorithm (beyond Splay)
- ▶ made several existing results more intuitive, in particular lower bounds for *OPT*
- ▶ but has not fueled the resolution of the Splay tree conjecture (yet?)

Arborally Satisfied Point Sets

Each proper rectangle must contain another point.

Definition 2.17 (Arborally Satisfied)

A 2D point set X is arborally satisfied if it does not contain any unsatisfied rectangles.

An *unsatisfied rectangle* is formed by $p_1, p_2 \in X$, $p_1 = (x_1, y_1)$, $p_2 = (x_2, y_2)$ with

- ▶ with $x_1 \neq x_2$ and $y_1 \neq y_2$ and
- ▶ $X \cap [x_1, x_2] \times [y_1, y_2] = \{p_1, p_2\}$



BST-trace interpretation: $X_t := X \cap \mathbb{R} \times \{t\}$ corresponds to BST nodes touched at time t

BST $\iff$ Arborally Satisfied

Theorem 2.18

A point set $Y \subset [n] \times [m]$ corresponds to valid BST algorithm *iff* Y is arborally satisfied. $\blacktriangleleft$

Minimum Arborally Satisfied Superset

Given: Set of points $X \subset [n] \times \mathbb{N}$

Goal: smallest **arborally satisfied** superset $Y \supseteq X$

- ▶ care particularly for $X = \{(x_t, t) : t \in [m]\}$ for accesses $x_1, \dots, x_m$
- ▶ cost of BST algorithm corresponding to Y : $|Y|$
- ▶ If we're given all of X up front $\rightsquigarrow$ offline (BST) algorithm
- ▶ If X is revealed one time-slice X_t at a time $\rightsquigarrow$ online algorithm for geometric view

Greedy Arborally Satisfied Superset

Natural Greedy Algorithm: GreedyASS

- ▶ at time t , add the points now (current time row) that are necessary to satisfy $Y \cup X_t$

↪ geometric sweep-line algorithm

- ▶ Seems to work very well
- ▶ but easy to see that it has approximation ratio $\geq \frac{3}{2}$

Who controls the Future

Amazing fact: GreedyASS translated to BSTs is GreedyFuture! 🐱

- ▶ to make the treap-construction effective, we need the future accesses (!) . . .
- ▶ but can delay actual treapify until next access comes
- ▶ can simulate this in BST (“*split trees*”) in constant amortized time

can be implemented as

- ▶ GreedyASS is an **online** BST algorithm

↪ New contender for a dynamically optimal algorithm!

👍 Indeed, by now more positive results known about Greedy than Splay

👎 probably mostly of theoretical interest as a BST algorithm itself

Lower Bounds

Geometric view also yields lower bounds on $OPT(X) = \min\{|Y| : Y \supset X, Y \text{ a.s.s.}\}$

- Independent rectangles (in X): no rectangle contains other rectangle's **corner** in its **inside**

Theorem 2.19

$OPT(X) \geq |X| + \frac{1}{2} \max |I(X)|$ with max over set of *independent rectangles* $I(X)$.



Independent Rectangle Approximation

Approximate lower bound

- ▶ Computing $\max |I(X)|$ is not easy in general
- ▶ proof shows that signed independent rectangles suffice
- ↪ Signed Greedy: satisfy only positive resp. only negative independent rectangles
- ▶ two valid lower bounds, use larger one!
- ▶ Can prove: Signed Greedy gives constant factor approximation to best independent-rectangle bound
- ≠ constant factor within *OPT* (at least not known!)

2.9 Deferred Data Structures

Motivation

All our adaptive dictionaries so far were binary search trees.

↪ *We may be able to structure them to have cheap access to “important” items . . .*

But we inevitably keep all elements of the dictionary perfectly sorted (BST!)

↪ Most insertions must cost $\Omega(\log n)$

- ▶ Reduction from sorting

- ▶ BST-Sort: Insert all n items & output in inorder traversal

How can you possibly do better in a comparison-based model?

- ▶ Seems that we need sortedness for any fast queries, so sorting inevitable

↪ True. What if we never query (certain elements), though?

- ▶ Insertion-heavy workloads rather common scenario in applications

- ▶ many insertions, few queries ↪ most elements never queried for! *So why sort them?*

The Multiple Selection Problem

► Selection by Rank:

- Given: n unsorted elements $x_1, \dots, x_n$,
- Goal: find the r th smallest $x_{(r)}$ (at (1-based) index r after sorting)

► Multiple Selection:

- Given: n unsorted elements $x_1, \dots, x_n$,
- Goal: find q elements of ranks $r_1 < r_2 < \dots < r_q$ from unsorted elements $x_1, \dots, x_n$
 $\rightsquigarrow$ Report sought elements $x_{(r_1)}, \dots, x_{(r_q)}$ in sorted order

► Example:

x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8	x_9	x_{10}	x_{11}	x_{12}	x_{13}	x_{14}	x_{15}
67	30	45	33	15	99	26	90	55	9	96	45	95	31	3

Answer: 3, 9, 15, 45

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
3	9	15	26	30	31	33	45	45	55	67	90	95	96	99
↑	↑	↑					↑							
r_1	r_2	r_3					r_4							
1	2	3					8							

Recall: The Median-of-Medians Algorithm

```


1 procedure choosePivotMoM( $A[l..r]$ ):
2    $m := \lfloor n/5 \rfloor$ 
3   for  $i := 0, \dots, m-1$ 
4      $\text{sort}(A[5i..5i+4])$ 
5     // collect median of 5
6      $\text{Swap } A[i] \text{ and } A[5i+2]$ 
7   return  $\text{quickselectMoM}(A[0..m], \lfloor \frac{m-1}{2} \rfloor)$ 
8
9 procedure quickselectMoM( $A[l..r], k$ ):
10  if  $r - l \leq 1$  then return  $A[l]$ 
11   $b := \text{choosePivotMoM}(A[l..r])$ 
12   $j := \text{partition}(A[l..r], b)$ 
13  if  $j == k$ 
14    return  $A[j]$ 
15  else if  $j < k$ 
16     $\text{quickselectMoM}(A[j+1..r], k)$ 
17  else //  $j > k$ 
18     $\text{quickselectMoM}(A[l..j], k)$ 

```

Analysis:

► Note: 2 mutually recursive procedures
 $\rightsquigarrow$ effectively 2 recursive calls!

1. recursive call inside choosePivotMoM on $m \leq \frac{n}{5}$ elements
2. recursive call inside quickselectMoM



$\rightsquigarrow$ partition excludes $\sim 3 \cdot \frac{m}{2} \sim \frac{3}{10}n$ elem.

$$\rightsquigarrow C(n) \leq \Theta(n) + C\left(\frac{1}{5}n\right) + C\left(\frac{7}{10}n\right)$$

$$\begin{aligned} &\leq \Theta(n) + C\left(\frac{1}{5}n + \frac{7}{10}n\right) \\ &= \Theta(n) + C\left(\frac{9}{10}n\right) \rightsquigarrow C(n) = \Theta(n) \end{aligned}$$

ansatz: overall cost linear

Multiple Selection – Multiple Solutions

Several algorithmic ideas possible:

1. q calls to selection algorithm (quickselect, median of medians) $\rightsquigarrow O(qn)$
clearly wasteful

2. **Divide & Conquer over Ranks:**

(single) select $r_{\lceil q/2 \rceil}$ -th smallest and partition $x_1, \dots, x_n$ around it.

Recursively select $r_1, \dots, r_{\lceil q/2 \rceil - 1}$ and $r_{\lceil q/2 \rceil + 1}, \dots, r_q$

$\rightsquigarrow O(n \lg q)$

3. **Divide & Conquer over Elements:**

Find **median of** $x_1, \dots, x_n$ and split query ranks.

Recurse on subproblem only of it contains query ranks.

$\rightsquigarrow O(n \lg q)$

Can also show: $\Omega(n \lg q)$ comparisons needed in worst case. $\rightsquigarrow$ *Case closed?*


Multiple Selection – Lower Bound

... not if we put our adaptive-analysis glasses on!

“Gap Entropy”:

$$\mathcal{B} = \sum_{i=1}^{q+1} \Delta_i \log_2 \frac{n}{\Delta_i}$$

with $\Delta_i = r_i - r_{i-1}$ ($1 \leq i \leq q+1$, $r_0 = 0$ and $r_{q+1} = n$)



Theorem 2.20 (Multiple Selection Lower Bound)

Multiple selection requires $\geq \mathcal{B} - O(n)$ comparisons and $\Omega(\mathcal{B} + n)$ time in the worst and average case over inputs of n elements and query ranks $r_1 < \dots < r_q$. ◀

Multiple Selection – Lower Bound Proof

Proof:

Suppose we have multiple-selection algorithm $\mathcal{A}$.

If we run $\mathcal{A}$ using $r_1 < \dots < r_q$, we get output $x_{(r_1)}, \dots, x_{(r_q)}$.

For $\mathcal{A}$ to be correct, it needs to learn how every other element compares to $x_{(r_1)}, \dots, x_{(r_q)}$.

We can thus build a *sorting algorithm* for $x_1, \dots, x_n$ by sorting the $\Delta_i - 1$ elements between $x_{(r_{i-1})}$ and $x_{(r_i)}$, for $i = 1, \dots, q$.

Since sorting is subject to the $\lg(n!)$ lower bound, we must use $\geq \lg(n!)$ in total.

For $\mathcal{A}$'s comparisons, this means

$$\begin{aligned}\text{\#comparisons} &\geq \lg(n!) - \sum_{i=1}^{q+1} \lg((\Delta_i - 1)!) \\ &= n \lg n \pm O(n) - \sum_{i=1}^{q+1} (\Delta_i - 1) \lg(\Delta_i - 1) \pm O(\Delta_i)\end{aligned}$$

Multiple Selection – Lower Bound Proof [2]

Proof (cont.):

using $(x - 1) \lg(x - 1) = x \lg x \pm O(\lg x)$ as $x \rightarrow \infty$, we get

$$\begin{aligned} \text{\#comparisons} &\geq n \lg n - \sum_{i=1}^{q+1} \Delta_i \lg(\Delta_i) \pm O(\Delta_i + \lg \Delta_i) \pm O(n) \\ &= \sum_{i=1}^{q+1} \Delta_i \lg(n/\Delta_i) \pm O(n). \end{aligned}$$

An Adaptive Multiple-Selection Algorithm

Our last approach, **Divide & Conquer over Elements**, is indeed **optimal** (up to constant factors)

```
1 procedure MultiSelect( $X[1..n]$ ,  $R[1..q]$ ):  
2   if  $r == \emptyset$  return  
3    $p := \text{median}(X[1..n])$  // linear-time selection  
4    $X_1, X_2 := \text{partition}(X[1..n], p)$   
5    $r := |X_1| + 1$  // rank of  $p$  in  $X$   
6    $R_1, R_2 := \text{partition}(R[1..q], r)$   
7   MultiSelect( $X_1, R_1$ )  
8   if  $r \in R$  then report  $p$  end if  
9   MultiSelect( $X_2, R_2$ )
```

Theorem 2.21 (MultiSelect Analysis)

Divide & Conquer over Elements using linear-time median selection finds ranks $r_1 < \dots < r_q$ among n unsorted elements in $O(\mathcal{B} + n)$ time. 

Multiple Selection Algorithms

One can indeed get very close to lower bound

- ▶ $\mathcal{B} + o(\mathcal{B}) + O(n)$ *comparisons* suffice



Kaligosi, Mehlhorn, Munro & Sanders: *Towards Optimal Multiple Selection*, ICALP 2005

- ▶ Using **Divide & Conquer over Elements**
with pivot chosen as median of random $n^{3/4}$ elements works
(analysis by induction)
- ▶ A more complicated algorithm also achieves the same deterministically

Going online

*Note that MultiSelect's recursion and partitioning does **not** depend on sought ranks*

- ▶ We stop when a subproblem has no rank
- ▶ other than that, just “median your way through”

↪ Can easily **add** further ranks *later*!

Quicksortus interruptus

- ▶ as MultiSelect, but remember pivots used in the past
- ▶ start between closest past pivots, repeatedly partition around median

↪ Same partitioning steps as in offline MultiSelect, just “on a funny schedule”

↪ analysis of Theorem 2.21 remains valid
(with $\mathcal{B}$ for the final set of queried ranks)

Online Multiple Selection

Lemma 2.22

The cost of one new query r'_i in (r_i, r_{i+1}) , splitting gap i of size Δ_i into $\Delta'_i = r'_i - r_i$ and $\Delta'_{i+1} = r_{i+1} - r'_i$ costs

$$\Delta'_i \lg\left(\frac{\Delta_i}{\Delta'_i}\right) + \Delta'_{i+1} \lg\left(\frac{\Delta_i}{\Delta'_{i+1}}\right) \in [\lg \Delta_i, \Delta_i]$$



2.10 Lazy Search Trees

Going dynamic

So we know how to add (and answer) more *queried ranks* in an online fashion


How about maintaining a dynamic set subject to insertions and deletions?

- ▶ Maintaining data in single array inconvenient

↪ Linked lists work fine!

- ▶ key property: access is *scanning* based
 - ▶ partitioning can directly work with linked lists (even reusing nodes)
 - ▶ median-of-medians pivot selection can also directly work with scans
-
- ▶ Insertion can use pivots to find correct gap, add element to list

Wishful-Thinking Solution



- ▶ BST of unsorted “bags”
 - ▶ Insert simply finds right bag, adds element
 - ▶ To answer query must look inside bag
- ~> partition one bag around query

Is this good? What running times should we aim for? Can we do better?

Notation

We mostly use the notation from multiple selection, with a few convenience changes

- ▶ m = number of gaps
- ▶ Gaps $\Delta_1, \dots, \Delta_m$ partition elements into bags (previously Δ_i denoted only the size)

$$\sum_{i=1}^m |\Delta_i| = n$$

- ▶ queried elements are smallest or largest in gap
- ▶ $r_0 = 0, r_{q+1} = n \rightsquigarrow |\Delta_i| = r_i - r_{i-1}$

Working against $\mathcal{B}$

Idea: Any snapshot still a multiple-select problem, subject to lower bound $\Omega(\mathcal{B} + n)$

↪ Trace how $\mathcal{B} + n$ changes, try to match that as amortized running time.

Lemma 2.23 (Query lower bound)

Suppose a query splits a gap Δ_i into two gaps of size x and $c \cdot x$ for $c \geq 1$.

Then the bound $\mathcal{B} + n$ increases by $\Omega(x \log c)$.


Moreover, $\mathcal{B} = \Omega(m \log n)$ for $m \geq 2$, so a query has $\Omega(\log n)$ amortized costs. ◀

Lemma 2.24 (Insert lower bound)

Suppose we insert an element into gap Δ_i .

Then the bound $\mathcal{B} + n$ increases by $\Omega\left(\log\left(\frac{n}{|\Delta_i|}\right)\right)$. ◀

Step 1: Biased search tree over gaps



$r_1 < r_2 < \dots < r_q$ queries ranks

$|\Delta_i| = r_i - r_{i-1}$

($r_0 = 0, r_{q+1} = n$)

Gap Data Structure

► store in *biased search tree* (by size $|\Delta_i|$) $\rightsquigarrow$ access Δ_i in $O\left(\log \frac{W}{w_i}\right) = O\left(\log \frac{n}{|\Delta_i|}\right)$

✓ Insert does find and changeWeight


✓ Query finds gap and replace it by two new ones, but $\Omega(\log n)$ increase in lower bound

The Query Problem

Recall that recursive median splitting was good enough for query in online multiple selection when set of elements static.

However, with (adversarial) insertions, medians don't stay medians . . .

Must maintain further structure inside gaps: *intervals*



Smaller intervals $\rightsquigarrow$ faster queries

More intervals $\rightsquigarrow$ slower insertions

Non-Solutions

- ▶ Collect insertions in buffers?
 - ⚡ $O(x \log c)$ amortized cost for queries?
 - ⚡ unclear how to balance gap data structure without flushing buffers
- ▶ Interval sizes roughly doubling as we move away from boundary.
 - ⚡ unclear if maintainable upon query in allowed time

Our Solution: Lazy Search Trees

In short: all operations amortized to lower bound, except $+O(\log \log n)$ per insert

Theorem 2.25 (Lazy search tree runtimes)

For n the current number of elements, $\{\Delta_i\}$ as above and q the total number of queries, we do

- ▶ $\text{Construction}(S)$ in $O(n)$ worst-case time, where $|S| = n$.
- ▶ $\text{Insert}(x)$ in $O(\min(\log(n/|\Delta_i|) + \log \log |\Delta_i|, \log q))$ worst-case time, where $x \in \Delta_i$.
- ▶ $\text{RankBasedQuery}(r)$ in $O(x \log c + \log n)$ amortized time, resulting gaps have size cx and $x, c \geq 1$.
- ▶ $\text{Delete}(\text{ptr})$ in $O(\log n)$ worst-case time.
- ▶ $\text{ChangeKey}(\text{ptr}, x')$ in $O(\min(\log q, \log \log |\Delta_i|))$ worst-case time, where the element at ptr , is $x \in \Delta_i$ and x' lies to its closest query rank in Δ_i ; otherwise, takes $O(\log n)$ worst-case time.
- ▶ $\text{Split}(r)$ in time according to $\text{RankBasedQuery}(r)$.
- ▶ $\text{Merge}(T_1, T_2)$ (where $T_1 \leq T_2$) in $O(\log n)$ worst-case time.



Define $B = \sum_{i=1}^m |\Delta_i| \log_2(n/|\Delta_i|)$. Then over a series of insertions and queries with no duplicate queries, the total complexity is $O(B + \min(n \log \log n, n \log q))$. ◀

 Sandlund, Wild: *Lazy Search Trees*, FOCS 2020

$O(\log \log n)$ in insert removed via selectable heaps

 Sandlund, Zhang: *Selectable Heaps and Optimal Lazy Search Trees*, SODA 2022

Lazy Search Trees – Overview



$r_1 < r_2 < \dots < r_q$ queries ranks

$|\Delta_i| = r_i - r_{i-1}$

$(r_0 = 0, r_{q+1} = N)$

Three levels:

1. Gaps Δ_i (= bags)

- maintained in a *biased search tree*

✓ access Δ_i in
 $O(\log \frac{W}{w_i}) = O(\log \frac{n}{|\Delta_i|})$

2. Within Δ_i , maintain set of *intervals* for fast queries

- $O(\log |\Delta_i|)$ intervals

- stored in BBST

☹ $O(\log \log n)$ time to access interval

3. Each interval is stored as *unsorted list*.

Intervals

Interval $\mathcal{J}_{i,j}$: Unsorted list with values between two pivots

► **Left/Right Intervals:**

interval $\mathcal{J}_{i,j}$ in gap Δ_i is on the *left side* if the closest query rank is to the left
Symmetrically for right.

► **Outside of an interval:**

$\mathbf{o}(\mathcal{J}_{i,j}) = \# \text{elements } \underline{\text{outside}} \text{ of } \mathcal{J}_{i,j} = \min\{|\mathcal{J}_{i,1}| + \dots + |\mathcal{J}_{i,j-1}|, |\mathcal{J}_{i,j+1}| + \dots + |\mathcal{J}_{i,\ell_i}|\}$

► **Invariant (#Int):** In Δ_i , the number ℓ_i of intervals $\mathcal{J}_{i,j}$ is at most $\ell_i \leq 4 \log_2 |\Delta_i|$

► when at risk (upon query), restored via interval-merging rule (below)

► sizes of intervals not restricted $\rightsquigarrow$ insert does no maintenance

► **Merge Rule (M):** If $|\mathcal{J}_{i,j}| + |\mathcal{J}_{i,j+1}| \leq \mathbf{o}(\mathcal{J}_{i,j})$, merge $\mathcal{J}_{i,j}$ and $\mathcal{J}_{i,j+1}$.

► Only triggered upon query; not an invariant

► This rule only applies between intervals of the same type;
a left interval is never merged with a right interval.

(M) implies (#Int)

Lemma 2.26

Right after applying rule (M), Invariant (#Int) holds.



2.11 Lazy Search Tree Operations

Insert


- (1) Find correct gap Δ_i (in biased search tree)
- (2) Find correct interval $\mathcal{J}_{i,j}$ (in balanced search tree)
- (3) Insert new element and increment weight of Δ_i

— no maintenance —

Rank-Based Query


- (1) Find correct gap Δ_i (in biased search tree)
- (2) Invoke Rule (M) on Δ_i
- (3) Find correct interval $\mathcal{J}_{i,j}$ (in balanced search tree)
- (4) Split interval $\mathcal{J}_{i,j}$ 5 times (see below)
- (5) Split gap and distribute intervals
- (6) Invoke Rule (M) on Δ'_i and Δ'_{i+1}

Interval before query:



Query point

Intervals after query:



Rank-Based Queries

Rank-Based Queries are all of the following

- ▶ Find the rank of element x (rank is the result)
- ▶ Select the element of rank r (rank r)
- ▶ Get element for key x (rank of x)
- ▶ Successor / Predecessor for key x (rank of returned element)
- ▶ min / max (rank 1 resp. n)

Potential

The amortized analysis uses a potential function

$$\Phi = N_{01} + \sum_{i,j} \Phi_{i,j}$$

$$\Phi_{i,j} = 1 + \max\{|J_{i,j}| - \mathbf{o}(J_{i,j}), 0\}$$

$$N_{01} = \text{\#elements in 0-sided and 1-sided gaps}$$

Lemma 2.27

Applying Rule (M) to Δ_i has amortized cost $O(\ell_i)$.



Amortized Costs

full details spelled out in



Rysgaard, Wild: *Towards Lazy B-Trees*, MFCS 2025

Lemma 2.28

The amortized cost of $\text{Insert}(x)$ into Δ_i is $O\left(\log \frac{n}{|\Delta_i|} + \log \log n\right)$




Lemma 2.29

The amortized cost of $\text{RankBasedQuery}(r)$ is $O(x \log c + \log n)$



Amortized Costs – Query Analysis

Case 1.1: 2-sided $\Delta_i, \mathcal{J}_{i,j}$ on left side



Lazy Search Trees – Discussion

- 👍 Lazy search trees (nearly) optimally solve insertions and queries compared to multiple-selection lower bound $\Omega(\mathcal{B} + n)$
- 👍 Never perform worse than standard BST: $O(\log n)$ amortized per operation
- 👍 When using splay tree as gap data structure and interval data structures inherit the **same access lemma guarantees as splay trees**
- 👍 When used like a **priority queue** (only query for min)
Lazy Search Trees behave like an efficient priority queue!
 - ▶ Insert, DecreaseKey in $O(\log \log n)$ time
 - ▶ DeleteMin in $O(\log n)$ time
- ↪ *Lazy Search Trees smoothly interpolate between priority queue and BST but do so automatically, according to actual usage*
- 👎 Optimality corresponds to worst case of inserting all elements first, then doing all queries