

ADVANCED

overall tree
= binary tree of mini trees

mini trees

micro trees

actual nodes

$\frac{1}{7} \lg n$ nodes

$\lg n$ nodes

n nodes

DATA STRUCTURES

4

Persistence

Advanced Data Structures · Summer 2026

Prof. Dr. Sebastian Wild

4 Persistence

- 4.1 Time Travel!
- 4.2 Application: Planar Point Location
- 4.3 Elementary ideas
- 4.4 Partial Success: Fat Nodes
- 4.5 Bounded Degree Pointer Machine
- 4.6 Partial Persistence

4.1 Time Travel!

Ephemeral Data Structures

Standard mode for data structures: destructive updates.

- ▶ Changes to the data structure destroy previous version
- ↪ queries and updates are always w.r.t. latest version

Persistent data structures have non-destructive updates

- ↪ Can still access past versions
 - ▶ useful for auditing
 - ▶ allows clever algorithms (see below)
 - ▶ natural mode of operation for functional programming
- ▶ **Partial persistence:** queries to any version, updates only to most recent
 - ▶ version history still a path ↪ number versions
- ▶ **Full persistence:** both queries and updates to any version
 - ▶ versions form a tree
- ▶ not here: more general versions thinkable, but unclear whether efficiently supportable
 - ▶ confluent persistence, functional persistence
 - ▶ retroactive data structures (time travel with single timeline . . .)

4.2 Application: Planar Point Location

Planar Point Location

Before diving into implementing persistence, let's see an application

4.3 Elementary ideas

Simple tricks

Full and Partial Persistence is trivial . . . if you don't care about time and space

- ▶ every update copies entire data structure
- ▶ keep a dictionary of versions

How (much) can we do better?

To make discussion concrete, let's try to build

Partially Persistent Doubly Linked Lists and Partially Persistent BSTs

- ▶ Upon one update *operation*, we may have a sequence of update *steps*:
 - ▶ update the value stored in a node
 - ▶ update pred/next pointer resp. left/right child pointer of a node
 - ▶ allocate a new node

↪ DLL and BST have few update steps per update operation

- ▶ copying entire data structure very wasteful

Attempt 1: Copy-On-Write Updates

Treat all fields in a node as *immutable*.

Upon update step at node x :

1. Create a copy x' of x with the new field value
2. For every other node y with a pointer to x , $y.p = x$, recursively trigger an atomic update $y.p := x'$
 - ▶ We assume here that these predecessors y of x are known
 - ▶ For DLLs trivial as always pointer in both directions exist
 - ▶ For BSTs, we can remember the predecessors during traversals

For header/root nodes, maintain sorted dictionary of *versions*

This effectively copies path from root to changed node $\rightsquigarrow$ *path copying*

Copy-on-Write – Discussion

- 👍 CoW directly supports full persistence
- 👍 For queries, only slowdown comes from locating the root of the desired version $O(\log m)$ in addition to original time after m updates.
- 👍 For low-depth, acyclic data structures, space overhead moderate
- 👎 DLLs copies all
 - ⚡ worst case $\Omega(nm)$ for ephemeral data structure with n nodes after m updates

Attempt 2: Log all Updates

Full replay log of all updates

- ▶ always store original data structure and full log of updates
- ▶ for every read and update, replay updates from beginning

👍 $O(n + m)$ space (assuming we can encode update in $O(1)$ space)

👎 $\Omega(i \cdot T_u + T)$ time for operation in version i
 T_u time per ephemeral update, T time for last operation

Attempt 3: Hybrid of CoW and Log

- ▶ store snapshot for every k th version

- ▶ update log in between

⇒ recompute version i from previous snapshot + log

👍 tune-able trade-off between space and time

👎 worst case still $\Omega(\sqrt{m})$ blow-up in either space or query time

4.4 Partial Success: Fat Nodes

Pointer-Based Data Structures

Recall assumption on data structure (generic version)

- ▶ Upon one update *operation*, we may have a sequence of update *steps*:
 - ▶ allocate a new node
 - ▶ update an *information field* of a node
 - ▶ update a *pointer field* of a node

↪ encompasses most data structures without arrays

Fat Nodes

Within each node and for each field: keep a log of writes, sorted by version stamp

- ▶ log stored as BBST
- ▶ when reading a field: find latest update before current version $\rightsquigarrow O(\log m)$ time
- ▶ writing field: add log entry $\rightsquigarrow O(\log m)$ time

👍 $O(1)$ extra space per update step

👎 $O(\log m)$ -factor slow-down for operations

Can be better if guaranteed to see few updates to fixed field!

4.5 Bounded Degree Pointer Machine

Terminology

Bounded In-Degree

4.6 Partial Persistence

General Transformation: Node Copying



Driscoll, Sarnak, Sleator, Tarjan: *Making data structures persistent*, JCSS 1989

Full Persistence

*Technique can be extended to full persistence at same costs(!)
unfortunately becomes more complicated*



Driscoll, Sarnak, Sleator, Tarjan: *Making data structures persistent*, JCSS 1989

Despite good properties of general transformation,
room for improvement on specific data structures.

~> Sizable literature on optimizations.

Time Travel Kills Amortization